

The Refutation Gap: Certifying Both Halves of an Optimality Claim in AI-Assisted Program Synthesis

Anonymous Author(s)

Abstract

Synthesis pipelines increasingly claim not just that a program is correct, but that it is *optimal*. Such a claim has two halves with radically different verification stories. The upper bound—“a program of size m exists”—is witnessed by an artifact that can be re-executed, proved equivalent to its specification, and shipped with a machine-checked certificate. The lower bound—“no program of size $m - 1$ exists”—is, in current practice, an unaudited assertion by a SAT solver. We call this the *refutation gap*, and we argue it is systematic rather than incidental: in the literature on minimal XOR circuits, three decades of published size records carry no certificate of any kind, and neither did 121 optimality results we ourselves produced. We study the gap in a working pipeline that synthesizes minimal linear straight-line programs over $\text{GF}(2)$. On the construction side we report a three-layer stack—re-execution, symbolic equivalence over the whole input space, and self-contained Lean 4 certificates—and give four case studies in which a layer caught a real defect that code review did not, twice overruling an explicit natural-language argument. On the refutation side we close the gap: every decisive UNSAT answer emits a DRAT proof that is checked by an independent third-party checker. All 62 refutations verify, agreeing with the uncertified results in every case, at a median proof size of 0.14 MB and a checking cost that is a small multiple of solving. We also report a practical obstacle for anyone attempting this: the fastest solver in our pipeline emits, through a standard Python interface, proofs the checker rejects. Finally we argue that in AI-assisted research the *write-up* is an artifact requiring verification too, and describe an adversarial audit of our own claims against raw data that found a double-counting defect and a statistic whose reported value lay outside the interval its data supported.

1 Introduction

Work on AI for verifiable coding has concentrated, reasonably, on the question *is this program correct?* A growing class of systems asks something stronger. Program-synthesis and discovery pipelines increasingly report that a found object is not merely valid but *best*: the shortest circuit, the smallest network, the tightest bound. Optimality claims of this kind now appear routinely in AI-driven discovery results.

An optimality claim is really a conjunction of two statements, and they are not symmetric:

upper bound	“a program of size m exists”
lower bound	“no program of size $m - 1$ exists”

The first is witnessed. Someone hands you an object; you re-run it, you check it against the specification, you can wrap it in a proof term and have a kernel confirm it. The whole apparatus of verifiable coding applies directly.

The second is not witnessed by anything. In practice it is produced by running a SAT solver until it returns UNSAT, and reporting that. The solver is tens of thousands of lines of heavily optimized C performing inprocessing, clause deletion, and variable elimination; a silent bug in any of it yields a lower bound that is simply wrong, with no artifact anyone can inspect. The verification effort goes entirely into the half that was already easy.

We call this the **refutation gap**. Our claim is that it is systematic. In the specific literature we study—minimal XOR circuits for cryptographic diffusion layers—the published record for the AES MixColumns matrix has moved nine times since 1997, most recently in 2026, and *not one* of those results is accompanied by a lower bound, let alone a certificate. We surveyed the primary sources and found no published lower bound for that matrix or for any MDS diffusion matrix. When we built our own pipeline we reproduced the pattern without noticing: we shipped three independent verification layers on the construction side and 121 optimality results resting on bare solver output.

Contributions.

1. We name and quantify the refutation gap in a working synthesis pipeline (§4), and show it is the norm in the surrounding literature.
2. We report a three-layer certification stack for synthesized GF(2) circuits and four case studies where a layer caught a real defect that inspection missed (§3).
3. We close the gap with proof-carrying refutations and measure the cost: 62/62 verified, median proof 0.14 MB, checking a small multiple of solving (§5). We report a solver-interface obstacle that cost us the faster solver.
4. We argue the research write-up is itself an artifact requiring verification, and describe an adversarial audit of our own claims that found two defects invisible to code-level checking (§6).

We are explicit about what is *not* new. SAT-based decision procedures for this synthesis problem date to 2010 [5], and were applied to linear matrices by Stoffelen [11]. Our contribution is not the encoding; it is the observation that the refutation half of these claims goes uncertified, a measurement of what certifying it costs, and an argument that pipelines making optimality claims should be held to it.

2 Setting: minimal XOR circuits as a testbed

Given a matrix M over GF(2), the *shortest linear straight-line program* (SLP) problem asks for the minimum number of XOR gates computing every row of M from the input variables. It is NP-hard [1, 2], and it is the standard cost model for the diffusion layer of a block cipher, so it carries a maintained table of records that specialists have pushed since 1997.

Three properties make it an unusually clean testbed for our question.

Verification is exact and instantaneous. A candidate program is a list of index pairs. Executing it symbolically gives, for each signal, the set of inputs XORed to produce it; the program is correct iff every target row appears. There is no tolerance, no sampling, no statistical test.

Optimality is decidable and small instances are in reach. “Is there a program with k gates?” reduces to SAT, so ground truth exists for small matrices—which is exactly what lets us study certification of the lower bound rather than merely asserting it.

Published baselines pin the implementation. Our implementation of Paar’s greedy algorithm reproduces the published gate counts exactly on six instances (AES MixColumns 108, ANUBIS 121, CLEFIA M1 121, KHAZAD 488, WHIRLPOOL 481), and our Boyar–Peralta implementation reproduces the published 97 on AES MixColumns [9, 15]. Reproducing a published number is weak evidence of correctness, but failing to is strong evidence of a defect—and it caught one for us (§3).

We stress that the SLP problem is a *case study*. The refutation gap is a property of any pipeline that claims minimality via a decision procedure.

3 Certifying the construction side

Our pipeline attaches three checks to every synthesized circuit. They share no code with the search, and—importantly—no code with each other.

Layer 1: re-execution. The circuit is re-run from scratch over the bitmask abstraction, in which a signal is the set of inputs XORed to form it. Microseconds. This is the gate on recording any result at all: a search output that fails it is never written to disk.

Layer 2: symbolic equivalence. Layer 1 works on an abstraction, and shares its notion of semantics with the code that built the problem instance. So we discharge a strictly stronger obligation with an SMT solver [4]: $\forall x. \text{circuit}(x) = Mx$ over GF(2), encoded as pure propositional logic since XOR is native. UNSAT of the negation is a proof over the entire input space, not the abstraction. 0.1–1 s per circuit. This catches a class of error Layer 1 structurally cannot: a mistake in the abstraction itself, or in how the specification was constructed.

Layer 3: proof terms. Each circuit is emitted as a self-contained Lean 4 file—no Mathlib, no build system, no imports—stating `SLP.valid n targets prog = true`, which unfolds to well-formedness (every gate reads only earlier signals; no gate is a self-XOR) and completeness (every target linear form is realized). Signals are represented as

coefficient vectors so that signal equality *is* equality of linear forms, leaving no abstraction gap to argue about. Each file ends with `#print axioms`, so the trust base is reported rather than assumed. Self-containment is deliberate: a certificate that requires resolving a library version is a certificate most readers will never check.

3.1 What the layers actually caught

The case for layered verification is usually made in the abstract. We can make it concretely: four defects, each caught by a check that shared no code with the thing it checked, and none caught by reading.

(i) **A search bug that inflated results.** Our rewrite of Paar’s algorithm rescanned signals it had created within the same loop, inflating gate counts and occasionally emitting a self-XOR. Layer 1 rejected it immediately. Code review had not.

(ii) **A wrong problem instance.** We implemented the CLEFIA diffusion matrices as circulant. They are Hadamard. Every result on those instances was being computed on the wrong matrix while passing all correctness checks—the circuits were valid, they just solved the wrong problem. The error surfaced only against an *external* anchor: after correction, our Paar count moved from 134 to 121, matching the published value exactly. This is the failure mode internal verification cannot reach, and it argues that reproducing published numbers should be treated as a verification layer in its own right.

(iii) **A malformed certificate.** Our Lean emitter line-wrapped inside `(a, b)` tuples, producing files that were not valid Lean. A round-trip test—parse the emitted file, compare to the source program—caught it. The generated output had been visually inspected and looked fine.

(iv) **Prose losing to arithmetic, twice.** A source review concluded that one cipher’s specification applies its diffusion matrix by right-multiplication, implying we should transpose. The argument was clear and we acted on it. Numerical cross-validation against independently derived test vectors then *failed* on the transposed matrix and *passed* on the original. A second instance: a size trend we had written up as a headline finding turned out to be an artifact of a counting bug. In both cases a natural-language argument was confidently wrong and a mechanical check was right. We think this is the central practical lesson for pipelines where an LLM supplies the reasoning: the reasoning is worth having, and it is not evidence.

4 The refutation gap

Having built all of that, we had 121 optimality claims—each of the form “the minimum is m ”—whose lower-bound half rested on a solver returning UNSAT. Nothing checked it. All three layers above certify *circuits*; none of them says anything about the non-existence of a smaller one.

This was not an oversight peculiar to us. Surveying the primary sources for the XOR-count literature, we found:

- The published record for AES MixColumns has moved nine times ($108 \rightarrow 103 \rightarrow 97 \rightarrow 95 \rightarrow 94 \rightarrow 92 \rightarrow 91 \rightarrow 89 \rightarrow 88$) between 1997 and 2026 [8–10, 13]. No entry is accompanied by a lower bound.
- We found no published lower bound for AES MixColumns or for any MDS diffusion matrix, in the record papers or elsewhere.
- Existing exact work [5, 11] proves optimality for small instances but, as is standard for its era, reports solver verdicts rather than certificates.

The gap is structural: the incentive is to report a better *object*, and an object is self-evidencing. A refutation has nothing to show.

5 Closing it: proof-carrying refutations

We modify the optimality procedure so that an optimality result is a **pair** of independently checkable objects:

(verified circuit of size m , checked DRAT refutation at $k = m - 1$).

The descent from a verified upper bound is unchanged; every SAT answer yields a circuit that is re-verified, and the first UNSAT answer now emits a DRAT proof [14] that is checked by `drat-trim`—a separate, third-party implementation

Instances re-established	71
certified by an independently checked refutation	62
closed by a counting bound, no refutation needed	9
uncertified	0
Disagreements with the earlier uncertified results	0
DRAT proof size, median / max	0.14 MB / 15.5 MB
Solve time with proof logging, median / max	0.02 s / 4.3 s
Independent check time, median / max	0.09 s / 6.9 s
Check/solve ratio, median	4.3×
Check/solve ratio, largest instance	1.6×

Table 1: Proof-carrying optimality. Every refutation the pipeline relies on is now checked by a third-party checker, and every previously uncertified result survived. The median check/solve ratio is dominated by the checker’s fixed startup cost on tiny instances; on the largest instance, where the ratio is meaningful, checking costs 1.6× solving.




Figure 1: Certification is affordable at this scale. **(a)** DRAT proof size against formula size; certificates stay small in absolute terms. **(b)** Independent check time against solve time, with the $y = x$ reference. Points above the diagonal at the left are the checker’s fixed startup cost, not a scaling problem; the trend converges toward the diagonal as instances grow.

sharing no code with the solver that produced it. A result is reported as proved *only* if the checker returns VERIFIED; a bare solver verdict no longer suffices.

5.1 Results

We re-established every previously closed instance under this regime.

Two things are worth drawing out. First, **nothing broke**: all 62 refutations verified and every optimality value agreed with what the uncertified runs had reported. That is the expected outcome, and it is worth stating plainly because it is the outcome that makes certification look unnecessary. It is not: the point is that before this we could not have distinguished this world from one in which a subtle solver or encoding bug had corrupted some fraction of our lower bounds, and neither can any of the published results that report bare verdicts.

Second, the cost is small. At these sizes proof logging did not measurably slow solving, proofs are megabytes rather than gigabytes, and checking costs roughly 1.6× solving where the ratio is meaningful.

5.2 An obstacle worth reporting

The faster solver in our pipeline is CaDiCaL. Through the standard Python interface we use [7], the proof it emits is *rejected* by `drat-trim(s NOT VERIFIED, “no conflict”)`, including after appending an explicit terminating empty clause. Glucose 4.2’s proofs verify without incident, so we use Glucose for proof-carrying runs.

We report this because it is exactly the sort of friction that decides whether practitioners certify or not. The path of least resistance is the fast solver with an unchecked verdict; making certification the default requires that the proof-emitting path be as easy to reach as the fast one. At our instance sizes the substitution happened to be free—Glucose was in fact marginally faster—but that is a property of small instances, not a general guarantee.

6 Verifying the claims, not only the code

A pipeline can be fully certified at the artifact level and still produce a false paper. The numbers that reach the reader are computed by aggregation code that nothing verifies, and framed by prose that nothing checks.

We ran an adversarial audit of our own write-up: an LLM agent was given the draft and the raw run records and instructed to recompute every quantitative claim and report discrepancies, with an explicit brief that overstatement was the target. It found two defects that no amount of code-level verification would have surfaced.

A double-counting defect in aggregation. The script that produced our headline table computed a deduplication key and never used it, so instances appearing in more than one sweep were counted twice. The reported count of closed instances was inflated by 20%, and—because duplicates were not uniformly distributed—a monotone trend we had presented as a central finding disappeared entirely once corrected.

A censored statistic. We reported the fraction of instances on which a heuristic was exactly optimal, computed over instances the solver could close. But closability is not independent of the quantity being measured: a larger gap requires more descent steps, each a chance to exhaust the budget. Correcting for the censoring, the true rate over all attempted instances lies in an interval that *excludes* the value we had reported.

Neither defect is exotic, and neither is visible to any artifact-level check. Both are the kind of error that an author is structurally poorly placed to find in their own work. We suggest that pipelines which use LLMs to help produce research should also use them adversarially against the resulting claims, and that the audit trail belongs in the artifact.

7 Limitations

Scale. Certified results here reach $n \leq 9$ inputs; the surrounding literature works at $n = 32$ and 64 . Our own frontier measurements show the encoding becoming erratic at $n = 10$ and not answering at $n = 11$. Nothing here suggests certified optimality is within reach for the matrices practitioners care about; the contribution is that where optimality *is* claimed, the refutation can and should be certified.

The checker is trusted. `drat-trim` is a third party relative to the solver, which is the property we need, but it is not itself formally verified. A stronger pipeline would emit LRAT and check it with a verified checker such as `cake_lpr` [12]; we regard that as the correct end state and did not reach it.

Symmetry breaking. Our encoding includes constraints that are sound for a *minimum*-size program but exclude some non-minimal programs of the same size. The DRAT proof certifies the encoded formula, not the argument that the encoding faithfully captures the problem. That argument remains a natural-language one, which—per §3(iv)—is precisely the kind of thing we have learned to distrust.

Lean certificates. The Lean 4 certificates are generated and their semantics are validated by a shadow evaluator that re-implements the definitions and includes negative controls. **[STATUS: compile output and axiom lists to be inserted before submission.]**

8 Related work

Proof-carrying results in combinatorial mathematics are well established: the Pythagorean triples problem [6] and subsequent large SAT results ship DRAT proofs as a matter of course, and verified checkers now exist for them [12]. Our observation is that this discipline has not transferred to program-synthesis pipelines that make optimality claims, even though the logical situation is identical.

For the SLP problem specifically, the SAT decision procedure is due to Fuhs and Schneider-Kamp [5], applied to linear matrices by Stoffelen [11]; neither reports proof certificates, which was standard practice at the time. The heuristic line—Paar’s algorithm, Boyar–Peralta [3], and its descendants [9, 13, 15, 16]—reports upper bounds only. On the verification side, our layering follows standard translation-validation practice; the contribution is the empirical account of what each layer caught and the extension of certification to the refutation half.

9 Conclusion

Optimality claims are becoming common in AI-assisted synthesis, and half of every such claim is currently unverified by construction. The half that gets the verification effort is the half that was already self-evidencing. We measured the cost of closing the gap in a real pipeline and found it small: proofs of megabytes, checking at a small multiple of solving, and no result disturbed. The obstacle is not cost but defaults—the fast path emits no proof, and the proof-emitting path is a step off the beaten track. We think pipelines that report minimality should be expected to ship the refutation, not just the object.

References

- [1] J. Boyar, P. Matthews, and R. Peralta. On the shortest linear straight-line program for computing linear forms. *MFCS*, LNCS 5162:168–179, 2008.
- [2] J. Boyar, P. Matthews, and R. Peralta. Logic minimization techniques with applications to cryptology. *Journal of Cryptology*, 26(2):280–312, 2013.
- [3] J. Boyar and R. Peralta. A new combinational logic minimization technique with applications to cryptology. *SEA*, 2010.
- [4] L. de Moura and N. Bjørner. Z3: An efficient SMT solver. *TACAS*, 2008.
- [5] C. Fuhs and P. Schneider-Kamp. Synthesizing shortest linear straight-line programs over $\text{GF}(2)$ using SAT. *SAT*, LNCS 6175:71–84, 2010.
- [6] M. J. H. Heule, O. Kullmann, and V. W. Marek. Solving and verifying the Boolean Pythagorean triples problem via cube-and-conquer. *SAT*, 2016.
- [7] A. Ignatiev, A. Morgado, and J. Marques-Silva. PySAT: A Python toolkit for prototyping with SAT oracles. *SAT*, 2018.
- [8] J. Jean. 88-XOR implementation of the AES MixColumns matrix. *IACR Cryptology ePrint Archive*, 2026/1481, 2026.
- [9] T. Kranz, G. Leander, K. Stoffelen, and F. Wiemer. Shorter linear straight-line programs for MDS matrices. *IACR Trans. Symmetric Cryptology*, 2017(4):188–211, 2017.
- [10] A. Maximov. AES MixColumn with 92 XOR gates. *IACR Cryptology ePrint Archive*, 2019/833, 2019.
- [11] K. Stoffelen. Optimizing S-box implementations for several criteria using SAT solvers. *FSE*, 2016.
- [12] Y. K. Tan, M. J. H. Heule, and M. O. Myreen. `cake_lpr`: Verified propagation redundancy checking in CakeML. *TACAS*, 2021.
- [13] Q. Q. Tan and T. Peyrin. Improved heuristics for short linear programs. *IACR Trans. CHES*, 2020(1):203–230, 2020.
- [14] N. Wetzler, M. J. H. Heule, and W. A. Hunt Jr. DRAT-trim: Efficient checking and trimming using expressive clausal proofs. *SAT*, 2014.
- [15] Z. Xiang, X. Zeng, D. Lin, Z. Bao, and S. Zhang. Optimizing implementations of linear layers. *IACR Trans. Symmetric Cryptology*, 2020(2), 2020.
- [16] Y. Yuan, W. Wu, T. Shi, L. Zhang, and Y. Zhang. A framework to improve the implementations of linear layers. *IACR Trans. Symmetric Cryptology*, 2024(2):322–347, 2024.